Queries, keys, and values

$$\{\boldsymbol{q}_i\}_{i=1}^{t} \rightsquigarrow \boldsymbol{Q} \in \mathbb{R}^{d' \times t}$$

$$\boldsymbol{q} = \boldsymbol{W}_{\boldsymbol{q}}\boldsymbol{x}, \quad \boldsymbol{k} = \boldsymbol{W}_{\boldsymbol{k}}\overset{\boldsymbol{\xi}}{\not{x}}, \quad \boldsymbol{v} = \boldsymbol{W}_{\boldsymbol{v}}\overset{\boldsymbol{\xi}}{\not{x}} \qquad \beta = \frac{1}{\sqrt{d'}}$$

$$\boldsymbol{q}, \boldsymbol{k} \in \mathbb{R}^{d'}, \quad \boldsymbol{v} \in \mathbb{R}^{d''}$$

$$\{\boldsymbol{\xi}_j\}_{j=1}^{\tau} \rightsquigarrow \{\boldsymbol{k}_j\}_{j=1}^{\tau}, \{\boldsymbol{v}_j\}_{j=1}^{\tau} \rightsquigarrow \boldsymbol{K}, \boldsymbol{V} \in \mathbb{R}^{\{d', d''\} \times \tau}$$

$$\boldsymbol{a} = \mathrm{softargmax}_{\beta}(\boldsymbol{K}^{\top}\boldsymbol{q}) \in \mathbb{R}^{\tau} \qquad \boldsymbol{h} = \boldsymbol{V}\boldsymbol{a} \in \mathbb{R}^{d''}$$

$$\{\boldsymbol{q}_i\}_{i=1}^{t} \rightsquigarrow \{\boldsymbol{a}_i\}_{i=1}^{t} \rightsquigarrow \boldsymbol{A} \in \mathbb{R}^{\tau \times t} \qquad \boldsymbol{H} = \boldsymbol{V}\boldsymbol{A} \in \mathbb{R}^{d'' \times t}$$

Self attention

$$d' = d'' \stackrel{\downarrow}{=} d$$

Implementation

$$\begin{bmatrix} \boldsymbol{q} \\ \boldsymbol{k} \\ \boldsymbol{v} \end{bmatrix} = \begin{bmatrix} \boldsymbol{W_q} \\ \boldsymbol{W_k} \\ \boldsymbol{W_v} \end{bmatrix} \boldsymbol{x} \in \mathbb{R}^{3d}$$

from the *RNN* lecture

$$\boldsymbol{h}[t] = g\left(\boldsymbol{W_h} \begin{bmatrix} \boldsymbol{x}[t] \\ \boldsymbol{h}[t-1] \end{bmatrix} + \boldsymbol{b_h}\right)$$

$$\boldsymbol{h}[0] \doteq \mathbf{0}, \boldsymbol{W_h} \doteq \begin{bmatrix} \boldsymbol{W_{hx}} \ \boldsymbol{W_{hh}} \end{bmatrix}$$

considering h heads we get a vector in $\mathbb{R}^{3hd}$

using a $\boldsymbol{W_h} \in \mathbb{R}^{d \times hd}$ to go back to $\mathbb{R}^d$

$$\begin{bmatrix} \boldsymbol{q}^1 \\ \boldsymbol{q}^2 \\ \vdots \\ \boldsymbol{q}^h \end{bmatrix} = \begin{bmatrix} \boldsymbol{W}_q^1 \\ \boldsymbol{W}_q^2 \\ \vdots \\ \boldsymbol{W}_q^h \end{bmatrix} \boldsymbol{x} \quad \begin{bmatrix} \boldsymbol{k}^1 \\ \boldsymbol{k}^2 \\ \vdots \\ \boldsymbol{k}^h \end{bmatrix} = \begin{bmatrix} \boldsymbol{W}_k^1 \\ \boldsymbol{W}_k^2 \\ \vdots \\ \boldsymbol{W}_k^h \end{bmatrix} \boldsymbol{x} \quad \begin{bmatrix} \boldsymbol{v}^1 \\ \boldsymbol{v}^2 \\ \vdots \\ \boldsymbol{v}^h \end{bmatrix} = \begin{bmatrix} \boldsymbol{W}_v^1 \\ \boldsymbol{W}_v^2 \\ \vdots \\ \boldsymbol{W}_v^h \end{bmatrix} \boldsymbol{x}$$

Transformer

Encoders-predictor-decoder architecture

(for Neural Machine Translation)

Transformer

x source sentence

y target sentence

$\tilde{y}$ predicted sentence

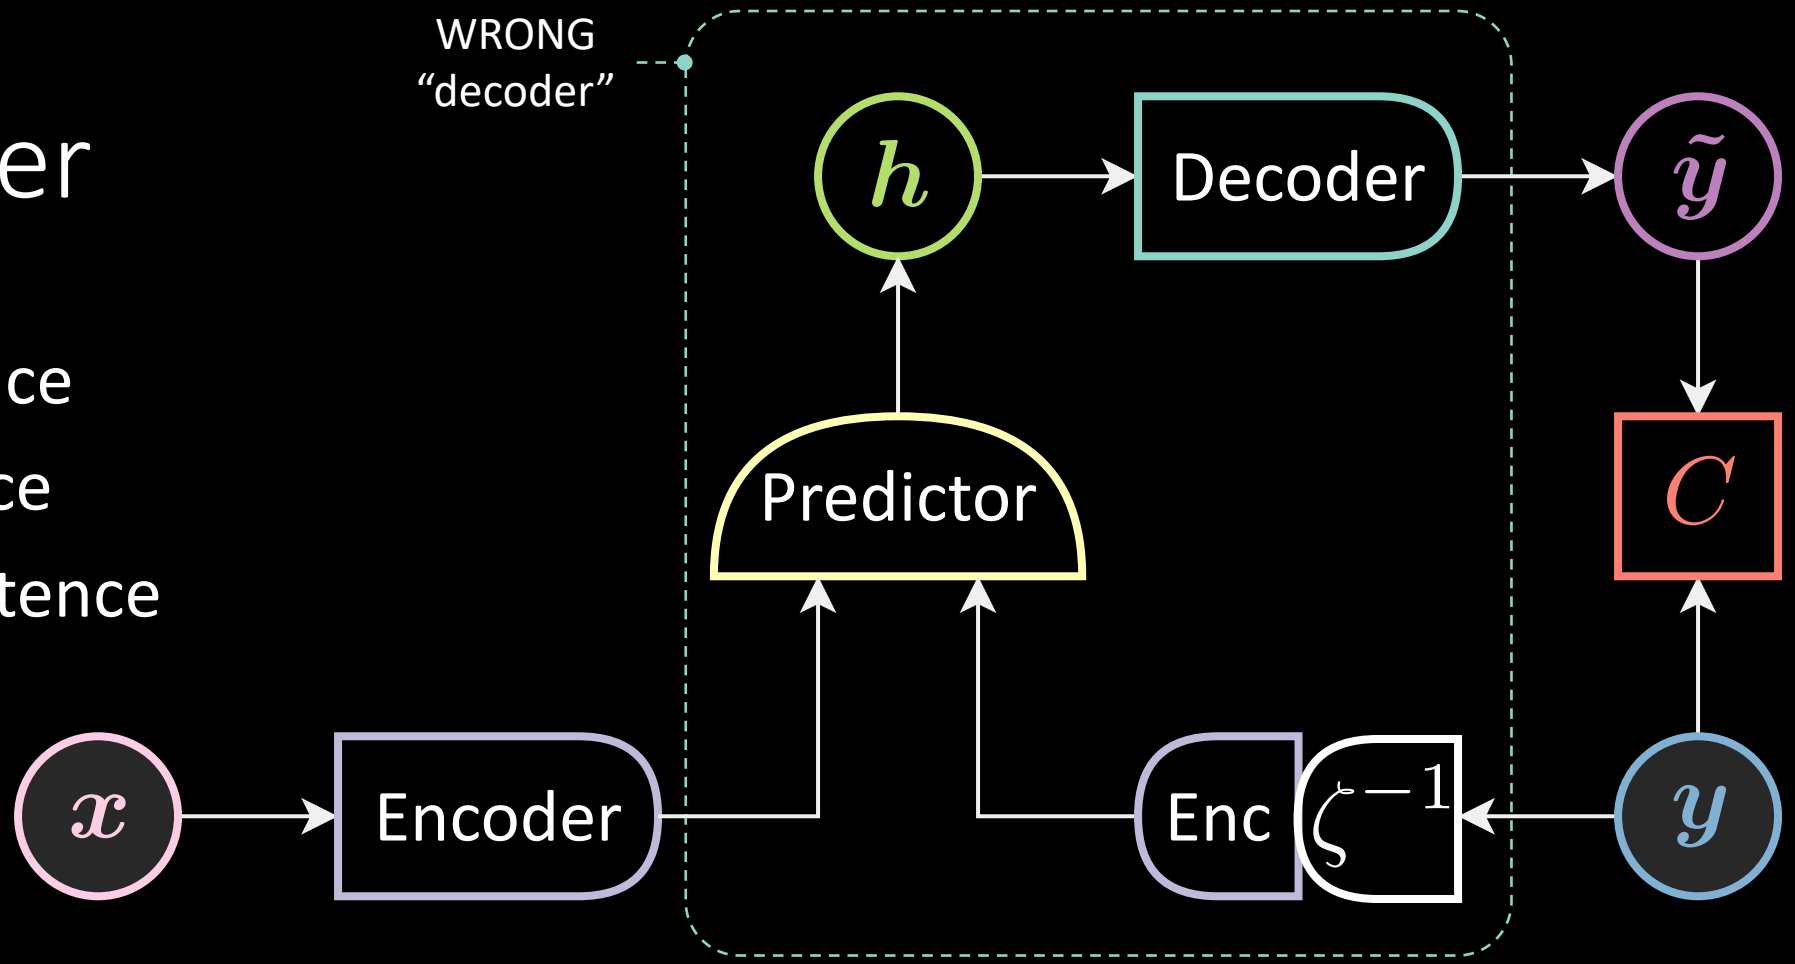

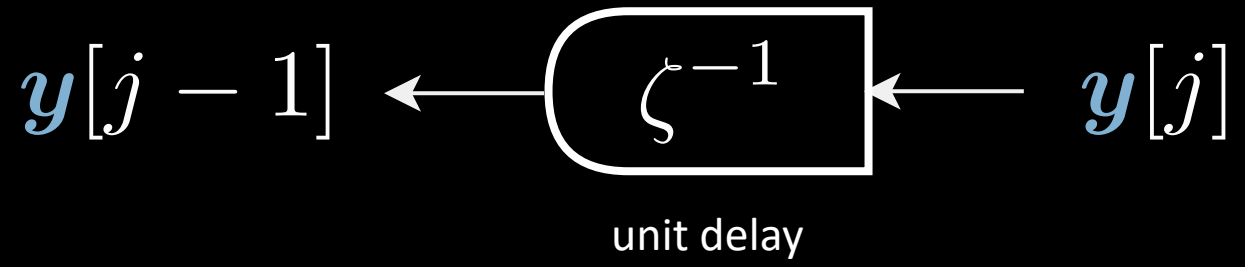

Transformer encoder

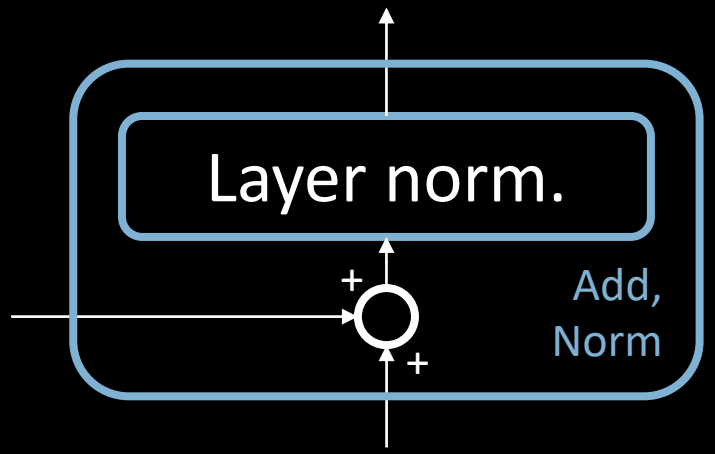

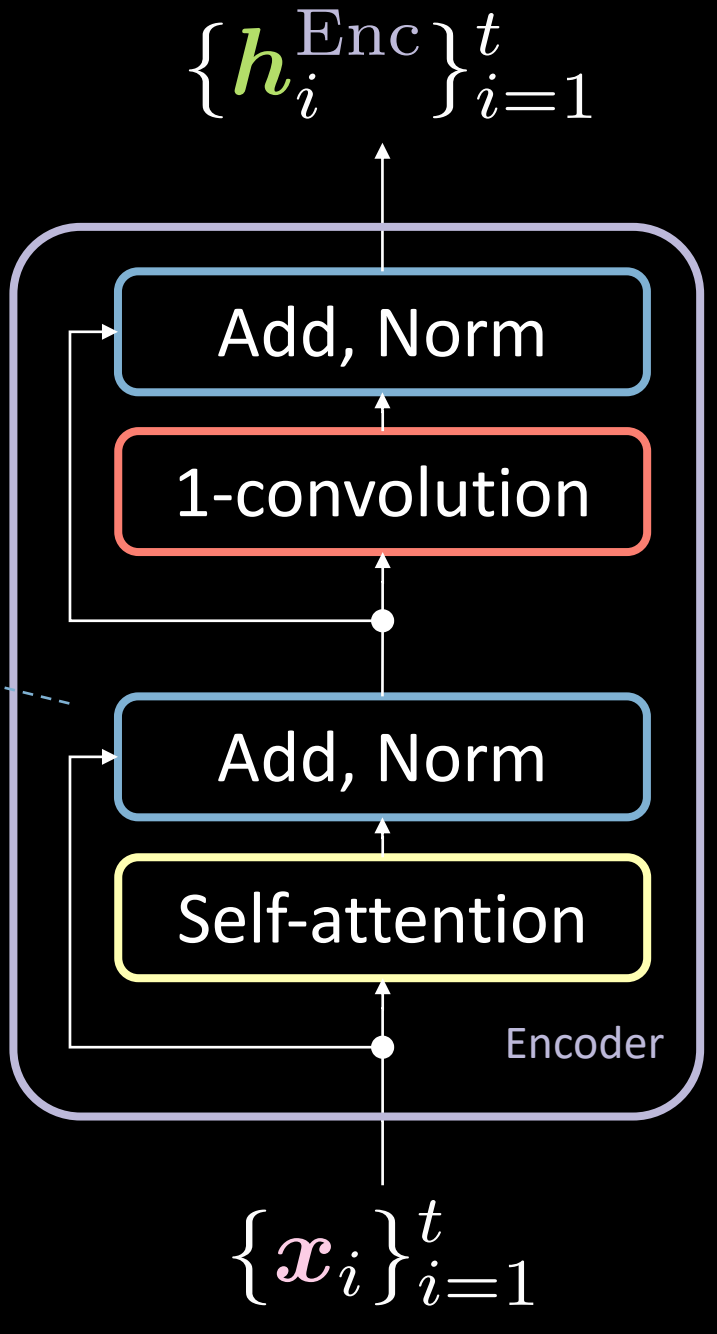

Transformer “decoder”

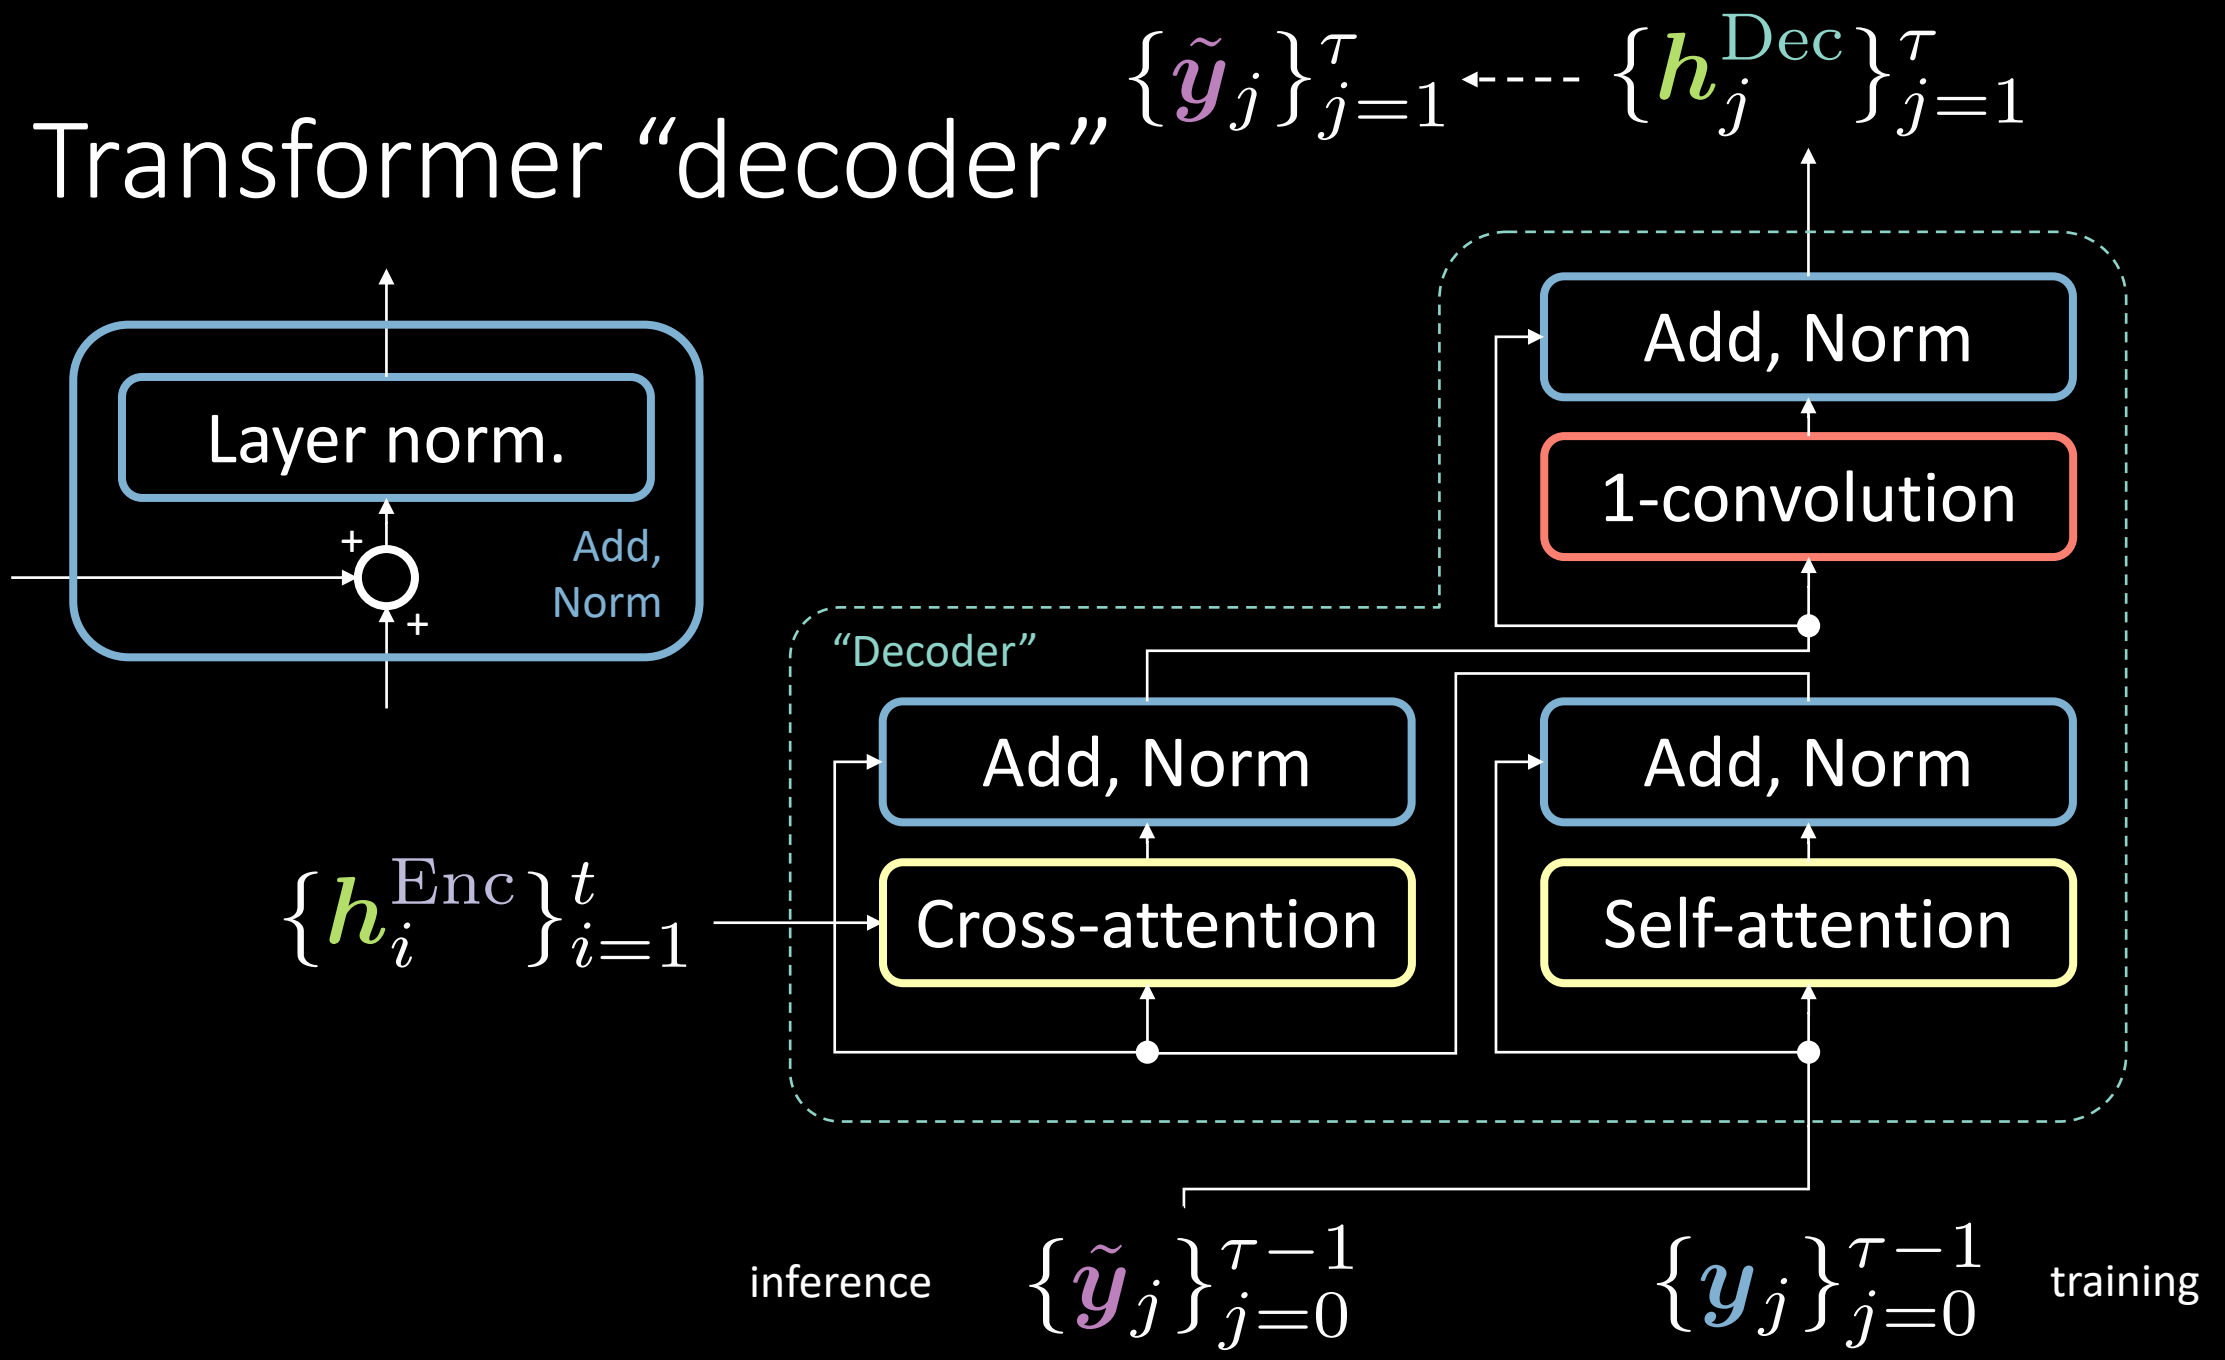

Transformer “decoder”